



# CS 540 Introduction to Artificial Intelligence

## **ML Intro / Unsupervised Learning I**

University of Wisconsin-Madison  
Fall 2023

# Artificial Intelligence

Machine learning

Deep learning with Artificial neural networks

Natural language processing

Computer vision

Robotics

Sept 26

Natural Language Processing (NLP)

Sept 28

Finish NLP; Machine Learning: Introduction

Oct 3

Machine Learning: Unsupervised Learning I

Oct 5

Machine Learning: Unsupervised Learning II

Oct 10

Machine Learning: Linear Regression

Oct 12

Machine Learning: K-Nearest Neighbors & Naive Bayes

# Outline


- Machine Learning Overview
  - Supervised learning, unsupervised learning, reinforcement learning
- Unsupervised Learning: Clustering
  - Hierarchical Clustering
    - Divisive, agglomerative, linkage strategies
  - Centroid-based, K-Means

# What is machine learning?


- Arthur Samuel (1959): the field of study that gives the computer the ability to learn **without being explicitly programmed**.
- Tom Mitchell (1997): A computer program is said to learn from **experience E** with respect to some class of **tasks T** and **performance measure P**, if its performance at tasks in T as measured by P, improves with experience E.




## Without Machine Learning



## With Machine Learning





# Taxonomy of ML



# Supervised Learning

## Supervised learning:

- Learn from labelled data.
- Dataset:  $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)$




- Goal: find function  $f : X \rightarrow Y$  to predict label on **new** data
- Labels can be discrete (“classification”) or real-valued (“regression”).

# Example 1: Predict whether a user likes a song or not



User Sharon






# Example 1: Predict whether a user likes a song or not



User Sharon

● Dislike

● Like




# Example 1: Predict whether a user likes a song or not




User Sharon

● Dislike


● Like



# Example 1: Predict whether a user likes a song or not






# Example 1: Predict whether a user likes a song or not



# Unsupervised Learning


- No labels; generally won't be making predictions
- Dataset:  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$
- Goal: find patterns/structures that help better understand data
  - E.g., dimension reduction, clustering, ...



Mulvey and Gingold


# Unsupervised Learning

Supervised Learning



# Clustering

- Given: dataset contains **no label**  $x_1, x_2, \dots, x_n$
- **Output:** divides the data into clusters such that there are intra-cluster similarity and inter-cluster dissimilarity



# Unsupervised Learning (UL)

- Clustering is just one type of unsupervised learning
  - PCA is another unsupervised algorithm
  - So is language modelling.
- Estimating probability distributions also UL (GANs)
- Clustering is popular & useful!



StyleGAN2 (Kerras et al '20)




# Reinforcement Learning



# Reinforcement Learning

- Given: an agent that can take actions and a reward function specifying how good an action is.
- **Goal:** learn to choose actions that maximize future reward total.



Google Deepmind

# Reinforcement Learning Key Problems

1. Problem: actions may have delayed effects.


— Requires **credit-assignment**

2. Problem: maximal reward action is unknown

— Exploration-exploitation trade-off

“..the problem [exploration-exploitation] was proposed [by British scientist] to be dropped over Germany so that German scientists could also waste their time on it.”

- Peter Whittle



Multi-armed Bandit

# Today: Clustering

- Several types of clustering

## Partitional

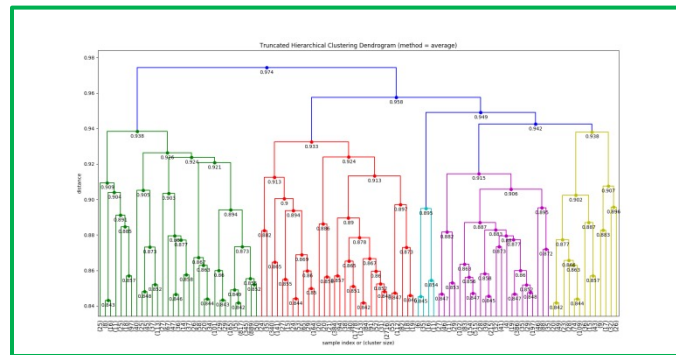

- Center-based
- Graph-theoretic
- Spectral

## Hierarchical

- Agglomerative
- Divisive

## Bayesian

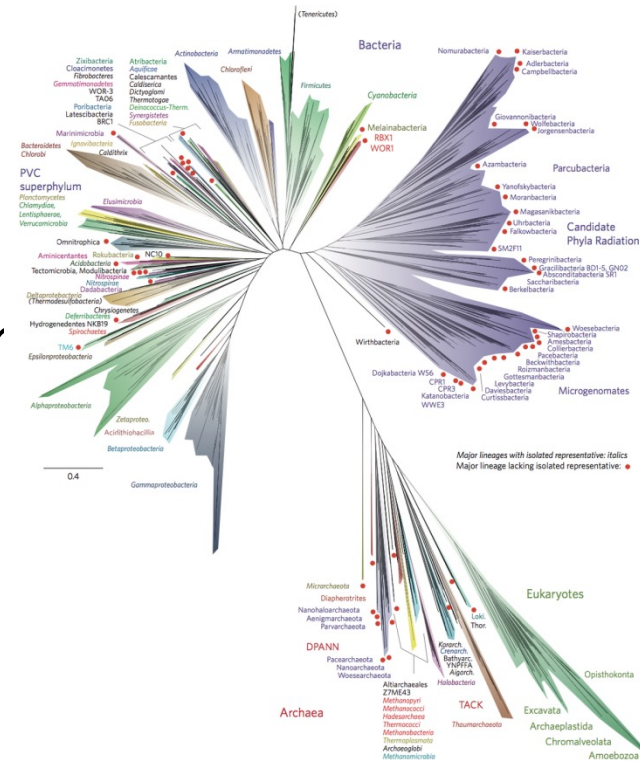
- Decision-based
- Nonparametric



# Hierarchical Clustering

Basic idea: build a “hierarchy”

- Want: arrangements from specific to general
- One advantage: no need for k, number of clusters.
- **Input:** points. **Output:** a hierarchy
  - A binary tree




Credit: Wikipedia

# Agglomerative vs Divisive

Two ways to go:

- **Agglomerative:** bottom up.
  - Start: each point a cluster. Progressively merge clusters
- **Divisive:** top down
  - Start: all points in one cluster. Progressively split clusters




# Hierarchical Agglomerative Clustering (HAC)

Input: data points  $x_1, \dots, x_n \in R^m$ , cluster distance function  $d(A, B)$

1. Initialize  $n$  clusters, one data point each
2. While (number of clusters  $> 1$ )
3.     find the closest clusters  $c_1, c_2 = \underset{A, B}{\operatorname{argmin}} d(A, B)$  over all cluster pairs  $A, B$
4.     merge  $c_1, c_2$  into a new cluster, remove  $c_1, c_2$

# Agglomerative Clustering Example


**Agglomerative.** Start: every point is its own cluster






# Agglomerative Clustering Example

**Get** pair of clusters that are closest and merge




# Agglomerative Clustering Example

**Repeat:** Get pair of clusters that are closest and merge



# Agglomerative Clustering Example

**Repeat:** Get pair of clusters that are closest and merge



# Cluster Distance Function

Merge: use closest clusters. Define closest?

- Single-linkage

$$d(A, B) = \min_{x_1 \in A, x_2 \in B} d(x_1, x_2)$$

- Complete-linkage

$$d(A, B) = \max_{x_1 \in A, x_2 \in B} d(x_1, x_2)$$


- Average-linkage

$$d(A, B) = \frac{1}{|A||B|} \sum_{x_1 \in A, x_2 \in B} d(x_1, x_2)$$

# Single-linkage Example

We'll merge using single-linkage

- 1-dimensional vectors.
- Initial: all points are clusters




# Single-linkage Example

We'll merge using single-linkage

$$d(C_1, \{4\}) = d(2, 4) = 2$$

$$d(\{4\}, \{5\}) = d(4, 5) = 1$$




# Single-linkage Example

Continue...


$$d(C_1, C_2) = d(2, 4) = 2$$

$$d(C_2, \{7.25\}) = d(5, 7.25) = 2.25$$




# Single-linkage Example

Continue...






# Single-linkage Example



# Complete-linkage Example

We'll merge using complete-linkage

- 1-dimensional vectors.
- Initial: all points are clusters




# Complete-linkage Example

Beginning is the same...


$$d(C_1, C_2) = d(1, 5) = 4$$

$$d(C_2, \{7.25\}) = d(4, 7.25) = 3.25$$




# Complete-linkage Example

Now we diverge:




# Complete-linkage Example



# When to Stop?

No simple answer:


- Use the binary tree (a **dendrogram**)
- Cut at different levels (get different heights/depths)



# Break & Quiz

**Q 1.1:** Let's do hierarchical clustering for two clusters with average linkage on the dataset below. What are the clusters?


- A.  $\{1\}, \{2, 4, 5, 7.25\}$
- B.  $\{1, 2\}, \{4, 5, 7.25\}$
- C.  $\{1, 2, 4\}, \{5, 7.25\}$
- D.  $\{1, 2, 4, 5\}, \{7.25\}$



# Break & Quiz

**Q 1.1:** Let's do hierarchical clustering for two clusters with average linkage on the dataset below. What are the clusters?

- A.  $\{1\}, \{2, 4, 5, 7.25\}$
- **B.  $\{1, 2\}, \{4, 5, 7.25\}$**
- C.  $\{1, 2, 4\}, \{5, 7.25\}$
- D.  $\{1, 2, 4, 5\}, \{7.25\}$





# Break & Quiz

**Q 1.2:** If we do hierarchical clustering on  $n$  points, the maximum depth of the resulting tree is

- A. 2
- B.  $\log n$
- C.  $n/2$
- D.  $n-1$

# Break & Quiz

**Q 1.2:** If we do hierarchical clustering on  $n$  points, the maximum depth of the resulting tree is


- A. 2
- B.  $\log n$
- C.  $n/2$
- **D.  $n-1$**

# Center-based Clustering

- k-means is an example of a partitional, **center-based clustering algorithm**.
- Specify a desired number of clusters,  $k$ ; run k-means to find  $k$  clusters.


# Center-based Clustering

- Steps: **1.** Randomly pick  $k$  cluster centers




# Center-based Clustering

- **2.** Find closest center for each point




# Center-based Clustering

- **3.** Update cluster centers by computing centroids



# Center-based Clustering

- Repeat Steps 2 & 3 until convergence



# K-means algorithm

- Input:  $x_1, x_2, \dots, x_n, k$
- Step 1: select  $k$  cluster centers  $c_1, c_2, \dots, c_k$
- Step 2: for each point  $x_i$ , assign it to the closest center in Euclidean distance:

$$y(x_i) = \operatorname{argmin}_j ||x_i - c_j||$$

- Step 3: update all cluster centers as the centroids:

$$c_j = \frac{\sum_{x:y(x)=j} x}{\sum_{x:y(x)=j} 1}$$

- Repeat Step 2 and 3 until cluster centers no longer change



# Break & Quiz

**Q 2.1:** You have seven 2-dimensional points. You run 3-means on it, with initial clusters

$$C_1 = \{(2, 2), (4, 4), (6, 6)\}, C_2 = \{(0, 4), (4, 0)\}, C_3 = \{(5, 5), (9, 9)\}$$

Cluster centroids are updated to?

- A.  $C_1: (4,4)$ ,  $C_2: (2,2)$ ,  $C_3: (7,7)$
- B.  $C_1: (6,6)$ ,  $C_2: (4,4)$ ,  $C_3: (9,9)$
- C.  $C_1: (2,2)$ ,  $C_2: (0,0)$ ,  $C_3: (5,5)$
- D.  $C_1: (2,6)$ ,  $C_2: (0,4)$ ,  $C_3: (5,9)$

# Break & Quiz

**Q 2.1:** You have seven 2-dimensional points. You run 3-means on it, with initial clusters

$$C_1 = \{(2, 2), (4, 4), (6, 6)\}, C_2 = \{(0, 4), (4, 0)\}, C_3 = \{(5, 5), (9, 9)\}$$

Cluster centroids are updated to?

- **A.  $C_1: (4,4)$ ,  $C_2: (2,2)$ ,  $C_3: (7,7)$**
- B.  $C_1: (6,6)$ ,  $C_2: (4,4)$ ,  $C_3: (9,9)$
- C.  $C_1: (2,2)$ ,  $C_2: (0,0)$ ,  $C_3: (5,5)$
- D.  $C_1: (2,6)$ ,  $C_2: (0,4)$ ,  $C_3: (5,9)$

# Break & Quiz

**Q 2.1:** You have seven 2-dimensional points. You run 3-means on it, with initial clusters

$$C_1 = \{(2, 2), (4, 4), (6, 6)\}, C_2 = \{(0, 4), (4, 0)\}, C_3 = \{(5, 5), (9, 9)\}$$

Cluster centroids are updated to?

- **A.  $C_1: (4,4)$ ,  $C_2: (2,2)$ ,  $C_3: (7,7)$**
- B.  $C_1: (6,6)$ ,  $C_2: (4,4)$ ,  $C_3: (9,9)$
- C.  $C_1: (2,2)$ ,  $C_2: (0,0)$ ,  $C_3: (5,5)$
- D.  $C_1: (2,6)$ ,  $C_2: (0,4)$ ,  $C_3: (5,9)$

The average of points in  $C_1$  is  $(4,4)$ .

The average of points in  $C_2$  is  $(2,2)$ .

The average of points in  $C_3$  is  $(7,7)$ .

# Break & Quiz

**Q 2.2:** We are running 3-means again. We have 3 centers,  $C_1$  (0,1),  $C_2$  (2,1),  $C_3$  (-1,2). Which cluster assignment is possible for the points (1,1) and (-1,1), respectively? Ties are broken arbitrarily:

(i)  $C_1, C_1$  (ii)  $C_2, C_3$  (iii)  $C_1, C_3$

- A. Only (i)
- B. Only (ii) and (iii)
- C. Only (i) and (iii)
- D. All of them

# Break & Quiz

**Q 2.2:** We are running 3-means again. We have 3 centers,  $C_1$  (0,1),  $C_2$  (2,1),  $C_3$  (-1,2). Which cluster assignment is possible for the points (1,1) and (-1,1), respectively? Ties are broken arbitrarily:

(i)  $C_1, C_1$  (ii)  $C_2, C_3$  (iii)  $C_1, C_3$

- A. Only (i)
- B. Only (ii) and (iii)
- C. Only (i) and (iii)
- **D. All of them**

# Break & Quiz

**Q 2.2:** We are running 3-means again. We have 3 centers,  $C_1$  (0,1),  $C_2$  (2,1),  $C_3$  (-1,2). Which cluster assignment is possible for the points (1,1) and (-1,1), respectively? Ties are broken arbitrarily:

(i)  $C_1, C_1$  (ii)  $C_2, C_3$  (iii)  $C_1, C_3$

- A. Only (i)
- B. Only (ii) and (iii)
- C. Only (i) and (iii)
- **D. All of them**

# Break & Quiz

**Q 2.2:** We are running 3-means again. We have 3 centers,  $C_1$  (0,1),  $C_2$  (2,1),  $C_3$  (-1,2). Which cluster assignment is possible for the points (1,1) and (-1,1), respectively? Ties are broken arbitrarily:

(i)  $C_1, C_1$  (ii)  $C_2, C_3$  (iii)  $C_1, C_3$

- A. Only (i)
- B. Only (ii) and (iii)
- C. Only (i) and (iii)
- **D. All of them**

For the point (1,1): square-Euclidean-distance to  $C_1$  is 1, to  $C_2$  is 1, to  $C_3$  is 5

So it can be assigned to  $C_1$  or  $C_2$

For the point (-1,1): square-Euclidean-distance to  $C_1$  is 1, to  $C_2$  is 9, to  $C_3$  is 1

So it can be assigned to  $C_1$  or  $C_3$

# Break & Quiz

**Q 2.3:** If we run K-means clustering twice with random starting cluster centers, are we guaranteed to get same clustering results? Does K-means always converge?

- A. Yes, Yes
- B. No, Yes
- C. Yes, No
- D. No, No



# Break & Quiz

**Q 2.3:** If we run K-means clustering twice with random starting cluster centers, are we guaranteed to get same clustering results? Does K-means always converge?

- A. Yes, Yes
- **B. No, Yes**
- C. Yes, No
- D. No, No

# Break & Quiz

**Q 2.3:** If we run K-means clustering twice with random starting cluster centers, are we guaranteed to get same clustering results? Does K-means always converge?

The clustering from k-means will depend on the initialization. Different initialization can lead to different outcomes.

- A. Yes, Yes
- **B. No, Yes**
- C. Yes, No
- D. No, No

K-means will always converge on a finite set of data points:

1. There are finite number of possible partitions of the points
2. The assignment and update steps of each iteration will only decrease the sum of the distances from points to their corresponding centers.
3. If it run forever without convergence, it will revisit the same partition, which is contradictory to item 2.